Control Flow

1 Learning Objectives

- Understand the use of condition codes and jump instructions in x86 assembly language.

- Recognize the components of simple loops in assembly language.

- Infer the C code corresponding to loops in assembly.

- Apply knowledge of the TEST and CMP instructions in the context of loops to trace switch statements in assembly.

2 Getting Started

To obtain a copy of today's activity, log into a shark machine and do the following:

```
$ wget http://www.cs.cmu.edu/~213/activities/machine-control.tar
$ tar xf machine-control.tar
$ cd machine-control
```

Record your answers to the discussion questions below. You may wish to refer back to the activity from the previous class (https://www.cs.cmu.edu/~213/activities/gdb-and-assembly.pdf) which contains a list of relevant GDB commands.

3 Basic Control Flow

This activity introduces the concept of condition codes and branch instructions.

The *condition codes* are four single-bit registers, named ZF, SF, CF, and OF. They are set implicitly by most arithmetic instructions (but not by MOV or LEA) and they have the following meanings:

ZF Result of operation was zero.

SF Result of operation was negative (its *Sign* bit was set)

CF Operation had an unsigned overflow (there was a *Carry* from the leftmost pair of bits)

OF Operation had a signed overflow (the sign bit of both inputs was the same, and the sign bit of the output is not equal to whatever that was)

Jump instructions (also called *branch instructions*) change the program counter (`%rip`). There are fifteen basic jump instructions:

Name	A.k.a.	Jump if...	After CMP...
JMP		Always	
JS		Negative (SF = 1)	
JNS		Not negative (SF = 0)	
JO		Signed overflow (OF = 1)	
JNO		No signed overflow (OF = 0)	
JE	JZ	Zero (ZF = 1)	Equal
JNE	JNZ	Not zero (ZF = 0)	Not equal
JB	JC, JNAE	Unsigned overflow (CF = 1)	Unsigned below
JAE	JNC, JNB	No unsigned overflow (CF = 0)	Unsigned above or equal
JA	JNBE	CF = 0 and ZF = 0	Unsigned above
JBE	JNA	CF = 1 or ZF = 1	Unsigned below or equal
JL	JNGE	SF ≠ OF	Signed less
JGE	JNL	SF = OF	Signed greater or equal
JG	JNLE	ZF = 0 and SF = OF	Signed greater
JLE	JNG	ZF = 1 or SF ≠ OF	Signed less or equal

The official "name" of a jump instruction is the name that `objdump` and `gdb` will use in disassembly listings. This is usually mnemonic for what the jump instruction will do if used immediately after a CMP instuction (as described in the "After CMP..." column). The "a.k.a." (also known as) names are mnemonic for other interpretations of what the instruction does; people writing assembly by hand can use them for clarity, but the distinction is lost in machine language.

Problem 1. Why is JZ (jump if zero) the same instruction as JE (jump if equal)?

"jump if ZF = 1" performs "jump if equal" when executed immediately after a CMP instruction.

Problem 2. Within the `machine-control` directory you created earlier, read the file `jumps.S`. The code in this file doesn't do anything *useful*, it just demonstrates the syntax of jump instructions. When you understand what's going on in this file, run these commands:

```
$ as jumps.S -o jumps.o
$ objdump -d jumps.o
```

Examine the output of the second command. (There will be a lot of output. You may want to make your shell window taller, or pipe the output to a "pager", e.g. `objdump -d jumps.o | less`). Compare it to what you remember from `jumps.S`, and the table above. You will probably notice that all of the "a.k.a." instructions have changed to the corresponding "name" instruction. What else do the lines for those groups of instructions have in common?

the first byte of the same instruction are the same

Problem 3. In the disassembly listing from the previous question, look at the *second* byte of each machine instruction. This is the part of the instruction that tells the CPU where to find the instruction that will be executed next (if the jump happens). Do you see a pattern to these bytes? What relationship is there among the address of "destination", the address of each jump instruction, the *length* of each jump instruction, and the value of the second byte?

those diff between two instructions is 2; the sec bytes add the address of next instructions = destination

Problem 4. *(Advanced)* Uncomment the line of `jumps.S` that reads

```
//    .skip  97
```

Repeat the `as` and `objdump` commands. What happened to the first several jump instructions? Why do you think this is? What would happen if you changed 97 to 98?

Experiment with various other numbers. Do you see a pattern?

the instruction has more bytes

Problem 5. *(Advanced)* Based on your answer to the previous question, what do you think the machine instruction `eb f0` would do?

4 Comparisons and Conditional Set Instructions

In this activity you will experiment with the CMP instruction, which sets the condition codes based on *comparing* two integers, and see how some of the conditional jump mnemonics correspond to some of C's relational operators. You will also be introduced to the *conditional set* instructions, which set a register to 0 or 1 based on the condition codes.

To begin this activity, run these commands (again, within the "machine-control" directory):

```
$ gdb ./cmp-set
(gdb) r
```

Read and follow the instructions that are printed, until it tells you to come back to this handout.

Problem 6. Based on the disassembly of `sete`, `seta`, and `setg`, which registers contain function arguments? Can you tell which is the first and which is the second argument?

di, si;NO

Problem 7. Based on the disassembly of `sete`, `seta`, and `setg`, which register contains the return value?

eax

Control Flow

You can call functions from the debugger with the `call` command.[1] For example,

```
(gdb) call sete(0, 1)
```

calls the function `sete` with arguments 0 and 1, and prints the result, like this:

```
(gdb) call sete(0, 1)
$1 = 0
```

In this case, `sete` returned zero. The "`$1 =`" prefix is to remind you that you can use `$1` in future function calls, or any other place GDB wants an arithmetic expression, to refer back to the number that was returned. (This is more useful with functions that can return many different values.)

Problem 8. Call `sete`, `seta`, and `setg` with each of the following pairs of values. Fill in the table.

Arg 1	Arg 2	sete	setg	seta
0	0	1	0	0
0	1	0	0	0
1	0	0	1	1
−1	0	0	0	1
0	−1	0	1	0
32768	32767	0	0	1
32767	32768	0	1	0
32768	32768	1	0	0
32768	−32768	1	0	0
−32768	32768	1	0	0
−32768	−32768	1	0	0

Problem 9. Assuming `%rdi` is the first and `%rsi` the second argument register, fill in the blanks in the C source code for `sete`, `setg`, and `seta`. (Hint: `stdint.h` defines the type name `int16_t` for 16-bit signed integers, and the type name `uint16_t` for 16-bit unsigned integers.)

```
#include <stdint.h>
int sete(uint16_t a,
         uint16_t b) {
    return a == b;
}
```

---

[1] **Caution:** Do not do this in bomb lab or attack lab. If you do, your bomb will explode, and your attack will not count.

```
int setg(   int16_t a,
            int16_t b) {
    return a >  b;
}

int seta(  uint16_t a,
           uint16_t b) {
    return a __ b;
}
```

5 > Tests and Conditional Move Instructions

In this activity you will experiment with the TEST instruction, which sets the condition codes based on the *bitwise and* of two integers. You will also be introduced to the *conditional move* instructions, which, based on the condition codes, either do or do not copy one register into another.

To begin this activity, run these commands (again, within the "machine-control" directory):

```
$ gdb ./test-cmov
(gdb) r
```

Read and follow the instructions that are printed, until it tells you to come back to this handout.

Problem 10. Why does the cmovc function contain a CMOVB instruction rather than a CMOVC instruction? (Hint: look carefully at the table on page 2.)

CMOVB and CMOVC are two names for the same machine instruction. The disassembler always prints "cmovb" for this instruction, even though it probably says "cmovc" in the source code for test-cmov.

Problem 11. In the disassembly of cmove, cmovs, and cmovc, what do you notice about the arguments to the TEST instruction?

they are the same

Problem 12. Call cmove, cmovs, and cmovc with each of the following pairs of values. Fill in the table.

Arg 1	Arg 2	cmove	cmovs	cmovc
0	0	0	0	0
0	1	1	0	0
0	2	2	0	0
1	0	0	0	0
1	1	0	0	0
1	2	0	0	0
−1	0	0	0	0
−1	1	0	1	0
−1	2	0	2	0
32767	1	0	0	0
32768	1	0	1	0

Problem 13. Why does cmovs(32768, 1) return 1?

32768 in 16bits is 0x8000 which is
less than 0

Problem 14. Is it possible to make cmovc return anything other than 0? Explain.

NO,never
carry

6 Loops

In Problem 5 we saw that jump instructions can jump both forward and backwards within the machine code. Backward jumps enable us to implement *loops*, in which part of the code is executed repeatedly.

Problem 15. You have been provided a file loops.o, containing machine code for three functions. The body of each function is a loop. Run the command

```
$ objdump -d loops.o
```

Translate the assembly language back into C and fill in the blanks in the functions below.

```c
int forLoop(int* x, int len) {
    int ret = 0;
    for (i = 0; i < len; i++) {
        ret += x[i];
    }
    return ret;
}

int whileLoop(int* x, int len) {
    int ret = 0;
    while (i < len) {
        ret += x[i];
        i++;
    }
    return ret;
}

int doWhileLoop(int* x, int len) {
    do {
        ret += x[i];
        i++;
    } while (i < len);
    return ret;
}
```

Problem 16. While you were working out the previous problem, how did you identify which register was used as the counter variable i?

One way is to look for a register that's being incremented by 1 each time around the loop. (Caution: we used -Og mode to compile loops.o. With more aggressive optimization, like what you get with -O2, there might not be any such register.

Problem 17. If we hadn't told you, and the names didn't give it away, could you have known that forLoop's C source contained a for loop and whileLoop's C source

contained a `while` loop?

NO;

7 Switch Statements

Switch statements in C are often compiled to *computed jumps* in assembly language. A jump instruction with an argument like

```
jmp     *.L4(,%rdi,8)
```

looks up the `%rdi`'th entry in the array beginning at `.L4`, and jumps to the address *stored in* that array entry. So, for instance, if `%rdi` is 2, and array entry 2 (counting from zero, as always) contains the address of label `.L5`, then the CPU will jump to `.L5`.

Here is a complete example of what this looks like in assembly.

```
switcher:
    cmpq $7, %rdi
    ja      .L2
    jmp     *.L4(, %rdi, 8)
.L7:
    xorq    $15, %rsi
    movq    %rsi, %rdx
.L3:
    leaq     112(%rdx), %rdi
    jmp     .L6
.L5:
    leaq    (%rdx, %rsi), %rdi
    salq     $2, %rdi
    jmp     .L6
.L2:
    movq     %rsi, %rdi
.L6:
```

```
    movq    %rdi, (%rcx)
    ret

    .section  .rodata
.L4:
    .quad         .L3
    .quad         .L2
    .quad         .L5
    .quad         .L2
    .quad         .L6
    .quad         .L7
    .quad         .L2
    .quad         .L5
```

Problem 18. The C code below is a partial translation ("decompilation") of the assembly code above. Fill in the case labels with the appropriate numbers.

```
// %rdi = a and val, %rsi = b, %rdx = c, %rcx = dest
void switcher(long a, long b, long c, long *dest) {
    long val;
    switch (a) {
case 5:
    c = b ^ 15;
case 0:
    val = c + 112;
    break;
case 7:
case 2:
    val = (c + b) << 2;
    break;
case 4:
    val = a;
    break;
default:
    val = b;
    }
    *dest = val;
}
```